

Predictive Modeling of Diabetes: A Machine Learning Approach

✧ Importing Dataset and Libraries

```

1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import graphviz
5 import matplotlib.pyplot as plt
6
7 !pip install ydata-profiling
8 from ydata_profiling import ProfileReport
9
10 from sklearn.preprocessing import LabelEncoder
11 from sklearn.preprocessing import MinMaxScaler
12 from sklearn.preprocessing import StandardScaler
13 from sklearn.feature_selection import RFE
14
15 from sklearn import tree
16 from sklearn.tree import plot_tree
17 from sklearn.tree import export_graphviz
18 from sklearn.model_selection import train_test_split
19 from sklearn.model_selection import StratifiedKFold
20 from sklearn.model_selection import cross_val_predict
21 from sklearn.tree import DecisionTreeClassifier
22 from sklearn.linear_model import LogisticRegression
23 from sklearn.neighbors import KNeighborsClassifier
24
25 from sklearn.metrics import confusion_matrix
26 from sklearn.metrics import classification_report
27 from sklearn.metrics import precision_score, accuracy_score, recall_score, f1_score

```

Preparing metadata (setup.py) ... done

Collecting phik<0.13,>=0.11.1 (from ydata-profiling)

Downloading phik-0.12.4-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (5.6 kB)

Requirement already satisfied: requests<3,>=2.24.0 in /usr/local/lib/python3.10/dist-packages (from ydata-profiling) (2.32.3)

Requirement already satisfied: tqdm<5,>=4.48.2 in /usr/local/lib/python3.10/dist-packages (from ydata-profiling) (4.66.6)

Requirement already satisfied: seaborn<0.14,>=0.10.1 in /usr/local/lib/python3.10/dist-packages (from ydata-profiling) (0.13.2)

Collecting multimethod<2,>=1.4 (from ydata-profiling)

Downloading multimethod-1.12-py3-none-any.whl.metadata (9.6 kB)

Requirement already satisfied: statsmodels<1,>=0.13.2 in /usr/local/lib/python3.10/dist-packages (from ydata-profiling) (0.14.4)

Requirement already satisfied: typeguard<5,>=3 in /usr/local/lib/python3.10/dist-packages (from ydata-profiling) (4.4.1)

Collecting imagehash==4.3.1 (from ydata-profiling)

Downloading ImageHash-4.3.1-py2.py3-none-any.whl.metadata (8.0 kB)

Requirement already satisfied: wordcloud>=1.9.3 in /usr/local/lib/python3.10/dist-packages (from ydata-profiling) (1.9.4)

Collecting dacite>=1.8 (from ydata-profiling)

Downloading dacite-1.8.1-py3-none-any.whl.metadata (15 kB)

Requirement already satisfied: numba<1,>=0.56.0 in /usr/local/lib/python3.10/dist-packages (from ydata-profiling) (0.60.0)

Collecting PyWavelets (from imagehash==4.3.1->ydata-profiling)

Downloading pywavelets-1.7.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (9.0 kB)

Requirement already satisfied: pillow in /usr/local/lib/python3.10/dist-packages (from imagehash==4.3.1->ydata-profiling) (11.0.0)

Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2<3.2,>=2.11.1->ydata-profiling) (3.0.2)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (1.3.1)

Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (4.55.0)

Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (1.4.7)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (24.2)

Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (3.2.0)

Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (2.8.2)

Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba<1,>=0.56.0->ydata-profiling) (0.43.0)

Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas!=1.4.0,<3,>1.1->ydata-profiling) (2024.2)

Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas!=1.4.0,<3,>1.1->ydata-profiling) (2024.2)

Requirement already satisfied: joblib>=0.14.1 in /usr/local/lib/python3.10/dist-packages (from phik<0.13,>=0.11.1->ydata-profiling) (1.4.2)

Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.10/dist-packages (from pydantic>=2->ydata-profiling) (0.7.0)

Requirement already satisfied: pydantic-core==2.23.4 in /usr/local/lib/python3.10/dist-packages (from pydantic>=2->ydata-profiling) (2.23.4)

Requirement already satisfied: typing-extensions>=4.6.1 in /usr/local/lib/python3.10/dist-packages (from pydantic>=2->ydata-profiling) (4.12.2)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.24.0->ydata-profiling) (3.4.0)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.24.0->ydata-profiling) (3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.24.0->ydata-profiling) (2.2.3)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.24.0->ydata-profiling) (2024.8.30)

Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.10/dist-packages (from statsmodels<1,>=0.13.2->ydata-profiling) (1.0.1)

12/2/24, 8:49 PM

Predictive Modeling of Diabetes using Decision Tree, K-NN and Logistic Regression - Colab

Downloading pywavelets-1.7.0-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (4.5 MB)

4.5/4.5 MB 77.9 MB/s eta 0:00:00

Building wheels for collected packages: htmlmin

Building wheel for htmlmin (setup.py) ... done

Created wheel for htmlmin: filename=htmlmin-0.1.12-py3-none-any.whl size=27081 sha256=5c1c3b118cfb088be8d90217d845f0cc6bc949a88fca5b2dfb0362b9b79e88

1 from google.colab import files

2 uploaded = files.upload()

Choose Files

Dataset of Diabetes .csv

Dataset of Diabetes .csv(text/csv) - 49511 bytes, last modified: 10/4/2024 - 100% done

Save Dataset of Diabetes .csv to Dataset of Diabetes .csv

1 db = pd.read_csv("Dataset of Diabetes .csv")

2 db

1000 rows × 14 columns

	ID	No_Patient	Gender	AGE	Urea	Cr	HbA1c	Chol	TG	HDL	LDL	VLDL	BMI	CLASS
0	502	17975	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
1	735	34221	M	26	4.5	62	4.9	3.7	1.4	1.1	2.1	0.6	23.0	N
2	420	47975	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
3	680	87656	F	50	4.7	46	4.9	4.2	0.9	2.4	1.4	0.5	24.0	N
4	504	34223	M	33	7.1	46	4.9	4.9	1.0	0.8	2.0	0.4	21.0	N
...
995	200	454317	M	71	11.0	97	7.0	7.5	1.7	1.2	1.8	0.6	30.0	Y
996	671	876534	M	31	3.0	60	12.3	4.1	2.2	0.7	2.4	15.4	37.2	Y
997	669	87654	M	30	7.1	81	6.7	4.1	1.1	1.2	2.4	8.1	27.4	Y
998	99	24004	M	38	5.8	59	6.7	5.3	2.0	1.6	2.9	14.0	40.5	Y
999	248	24054	M	54	5.0	67	6.9	3.8	1.7	1.1	3.0	0.7	33.0	Y

Next steps:

Generate code with db

View recommended plots

New interactive sheet

Stage 1: Data Preparation

Data Formatting

```
1 #Renaming the columns to remove capitalization, correct spelling and avoid programming constraints with the libraries being used. This allows for overall
2
3 db.columns = ['id', 'no_patient', 'gender', 'age', 'urea', 'cr',
4               'hba1c', 'chol', 'tg', 'hdl',
5               'ldl', 'vldl', 'bmi', 'db_class']
6
7 db.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   id           1000 non-null   int64
1   no_patient   1000 non-null   int64
2   gender       1000 non-null   object
3   age          1000 non-null   int64
4   urea         1000 non-null   float64
5   cr           1000 non-null   int64
6   hba1c        1000 non-null   float64
7   chol         1000 non-null   float64
8   tg           1000 non-null   float64
9   hdl          1000 non-null   float64
10  ldl          1000 non-null   float64
11  vldl         1000 non-null   float64
12  bmi          1000 non-null   float64
13  db_class     1000 non-null   object
dtypes: float64(8), int64(4), object(2)
memory usage: 109.5+ KB
```

Data Cleaning

```
1 #checking for missing values
2
```

```
3 print(db.isnull().sum())
```

```
id          0
no_patient  0
gender      0
age         0
urea        0
cr          0
hba1c       0
chol        0
tg          0
hdl         0
ldl         0
vldl        0
bmi         0
db_class    0
dtype: int64
```

```
1 #checking for duplicate values in the id column
2
3 duplicate_id = db['id'].duplicated().sum()
4 duplicate_id
```

```
200
```

```
1 #identifying the duplicate id vlaues
2
3 duplicate_num_id = db['id'].value_counts()
4 duplicate_num_id = duplicate_num_id[duplicate_num_id > 1]
5 print(duplicate_num_id)
```

```
id
76    2
108   2
57    2
26    2
69    2
..
150   2
49    2
144   2
145   2
147   2
Name: count, Length: 200, dtype: int64
```

```
1 #Random sample #1 - Checking if the ids are duplicate values for the id column or duplicate patient records (i.e. the both observations have the same val
2 #Since the other attributes have different values in each observation, this might be considered as a duplicate value for the id column only.
3
4 display(db[db['id'] == 76])
```

	id	no_patient	gender	age	urea	cr	hba1c	chol	tg	hdl	ldl	vldl	bmi	db_class
249	76	9903	M	73	4.3	79	6.0	5.3	1.4	1.5	3.2	0.6	27.0	Y
910	76	8978	M	60	5.4	64	10.4	3.8	1.5	0.8	2.3	0.6	31.0	Y

```
1 #Random sample #2 - Checking if the ids are duplicate values only or duplicate patient record (i.e. the both observations have the same values for each c
2 #Confirming that the duplicate ids can be consideres as duplicate values for id column only.
3
4 display(db[db['id'] == 150])
```

	id	no_patient	gender	age	urea	cr	hba1c	chol	tg	hdl	ldl	vldl	bmi	db_class
88	150	45382	F	40	6.3	79	4.9	4.3	0.8	0.8	1.8	1.2	22.0	N
184	150	34287	F	42	2.8	39	4.6	4.7	2.5	1.3	2.4	1.1	25.0	Y

```
1 #checking for duplicate values in the patient number column
2
3 duplicate_no_patient = db['no_patient'].duplicated().sum()
4 duplicate_no_patient
```

```
39
```

```
1 #identifying the duplicate patient numbers
2
3 duplicate_num_patient = db['no_patient'].value_counts()
4 duplicate_num_patient = duplicate_num_patient[duplicate_num_patient > 1]
5 print(duplicate_num_patient)
```

```
no_patient
454316    19
856        2
87654      2
71741      2
```

```

34290    2
14389    2
34517    2
48362    2
45646    2
44835    2
24033    2
34514    2
23972    2
34325    2
234      2
45370    2
34516    2
34518    2
2345     2
34515    2
34545    2
345      2
Name: count, dtype: int64

```

```

1 #Random sample #1 - Checking if the patient number are duplicate values only or duplicate patient record (i.e. the both observations have the same values
2 #Since the other columns have different values in each observation, this might be duplicate values for the patient number column only.
3
4 display(db[db['no_patient'] == 45646])

```



	id	no_patient	gender	age	urea	cr	hba1c	chol	tg	hdl	ldl	vldl	bmi	db_class
936	111	45646	F	63	4.7	55	12.60	7.9	5.0	1.3	2.1	0.7	35.0	Y
938	113	45646	F	55	2.1	23	9.96	4.1	4.2	1.2	1.4	1.3	29.0	Y

```

1 #Random sample #2 - Checking if the patient number are duplicate values only or duplicate patient record (i.e. the both observations have the same values
2 #Confirming that the duplicate patient numbers can be considered as duplicate values for the patient column only.
3
4 display(db[db['no_patient'] == 34290])

```



	id	no_patient	gender	age	urea	cr	hba1c	chol	tg	hdl	ldl	vldl	bmi	db_class
34	699	34290	F	47	5.6	67	5.1	6.5	1.5	0.9	4.9	0.7	23.0	N
260	510	34290	M	73	4.3	79	6.9	5.3	1.4	1.5	3.2	0.6	28.0	Y

```

1 #checking the gender column
2
3 gender_frequency = db['gender'].value_counts()
4 print(gender_frequency)

```



```

gender
M      565
F      434
f         1
Name: count, dtype: int64

```

```

1 #Converting the lowercase f to uppercase and updating the gender_frequency variable
2
3 db['gender'] = db['gender'].replace({'f': 'F'})
4 gender_frequency = db['gender'].value_counts()
5 print(gender_frequency)

```



```

gender
M      565
F      435
Name: count, dtype: int64

```

```

1 #checking the class column
2
3 class_frequency = db['db_class'].value_counts()
4 print(class_frequency)

```



```

db_class
Y      840
N      102
P       53
Y         4
N          1
Name: count, dtype: int64

```

```

1 #There is a space trailing after some of the Y and N values (eg: 'Y ' and 'N '). This is causing the data to be read as different values. Removing the sp
2
3 db['db_class'] = db['db_class'].replace({'N ': 'N', 'Y ': 'Y'})
4 class_frequency = db['db_class'].value_counts()
5 print(class_frequency)

```



```

db_class
Y      844

```

```
N    103
P     53
Name: count, dtype: int64
```

Stage 2: Exploratory Data Analysis

Statistical Summary

1 db.describe()

	id	no_patient	age	urea	cr	hba1c	chol	tg	hdl	ldl	vldl	
count	1000.000000	1.000000e+03	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	340.500000	2.705514e+05	53.528000	5.124743	68.943000	8.281160	4.862820	2.349610	1.204750	2.609790	1.854700	29.570000
std	240.397673	3.380758e+06	8.799241	2.935165	59.984747	2.534003	1.301738	1.401176	0.660414	1.115102	3.663599	4.960000
min	1.000000	1.230000e+02	20.000000	0.500000	6.000000	0.900000	0.000000	0.300000	0.200000	0.300000	0.100000	19.000000
25%	125.750000	2.406375e+04	51.000000	3.700000	48.000000	6.500000	4.000000	1.500000	0.900000	1.800000	0.700000	26.000000
50%	300.500000	3.439550e+04	55.000000	4.600000	60.000000	8.000000	4.800000	2.000000	1.100000	2.500000	0.900000	30.000000
75%	550.250000	4.538425e+04	59.000000	5.700000	73.000000	10.200000	5.600000	2.900000	1.300000	3.300000	1.500000	33.000000
max	800.000000	7.543566e+07	79.000000	38.900000	800.000000	16.000000	10.300000	13.800000	9.900000	9.900000	35.000000	47.700000

The average age of the patients in the dataset is 53.53, suggesting the population is mostly middle age. The creatinine variable has a large standard deviation (59.98) and a large range (from 6 to 800), suggesting the presence of extreme values.

Correlation Analysis

```
1 #Using Panda Profiling Report to perform analysis on the dataset. The output is hidden to ensure the complete download of the IPYNB file, and saving the
2 #When the output is showing the file cuts off at this code block
3
4 report = ProfileReport(db)
5 report
```

Show hidden output

Report confirms there are no missing values or duplicate rows within the dataset.

Both HbA1c and Cholesterol exhibit a normal bell shaped distribution, indicating a relatively even spread of data around the mean. The other features appear skewed, suggesting that their data distributions are more concentrated around certain values, and the presence of outliers.

Cholesterol, triglycerides, HDL and LDL have variances below 2, with the lowest being 0.44 for the HDL variable. This suggests that these variables exhibit small spread with most of the values being close to the mean.

There are several correlations observed among the variable. Triglycerides(tg) has a high positive correlation with vldl (0.599), and creatinine has a high positive correlation with urea (0.568). BMI exhibits moderately positive correlation with a coefficient of 0.417, suggesting elevated HbA1c levels are associated with BMI.

The target variable, db_class, has a strong positive correlation with bmi (0.549) and hba1c (0.662). This suggests that both bmi and hba1c are influential and important in determining the likelihood of diabetes.


There is a significant class imbalance with 844 individuals classified as diabetic, 103 as non-diabetic, and only 53 as pre-diabetic. This 51.5% imbalance may affect the performance of the predictive algorithms leading to bias towards the majority class.

Data Visualization

```
1 #Observing the values in Urea and Creatinine columns
2
3 plt.boxplot(db[['urea', 'cr']])
4 plt.title('Boxplot of Urea Levels and Creatinine Ratio Among Patients')
5
6 plt.xticks([1, 2], ['urea', 'cr'])
7 plt.ylabel("Values")
8
9 plt.show()
```



Boxplot of Urea Levels and Creatinine Ratio Among Patients




```

1 #Observing the values for the cholesterol and lipoproteins columns
2
3 plt.boxplot(db[['chol','tg','hdl','ldl','vldl']])
4 plt.title('Boxplot for Cholesterol and Lipoproteins Among Patients')
5
6 plt.xticks([1, 2, 3, 4, 5], ['chol','tg','hdl','ldl','vldl'])
7 plt.ylabel("Values")
8
9 plt.show()

```



Boxplot for Cholesterol and Lipoproteins Among Patients




There are noticeable outliers present in the cholesterol, lipoprotein, urea, and creatinine attributes. This suggests that several patients may have extreme values for these attributes.

Urea and creatinine ratio are clinical measurements used to assess kidney impairment. While cholesterol, the lipoproteins and triglycerides levels are used to determine cardiovascular damage. These outliers therefore suggest there may be patients with kidney and cardiovascular damage.

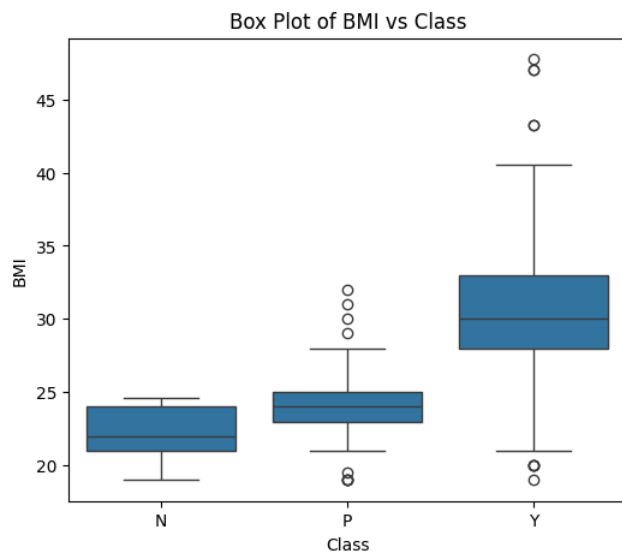
```

1 #Graph showing relationship between the BMI and Glucose attributes
2
3 plt.figure(figsize=(6, 5))
4 plt.scatter(db['bmi'], db['hba1c'], s=75, edgecolor='k')
5
6 plt.title('Scatter Plot of BMI vs HbA1c')
7 plt.xlabel('bmi')
8 plt.ylabel('hba1c')
9 plt.tight_layout()
10 plt.show()

```




```
1 #Graph showing relationship between the BMI and Class attributes
2
3 plt.figure(figsize=(6, 5))
4 sns.boxplot(x='db_class', y='bmi', data=db)
5 plt.title('Box Plot of BMI vs Class')
6 plt.xlabel('Class')
7 plt.ylabel('BMI')
8 plt.show()
```



```
1 #Graph showing relationship between the HbA1c and Class attributes
2
3 plt.figure(figsize=(6, 6))
4 sns.boxplot(x='db_class', y='hba1c', data=db)
5 plt.title('Box Plot of HbA1c vs Class')
6 plt.xlabel('Class')
7 plt.ylabel('HbA1c')
8 plt.show()
```



Box Plot of HbA1c vs Class



Stage 3: Data Preprocessing

Label Encoding

```

1 #Covertng the categorical attributes to numerical variables to make compatible for the classifiers.
2
3 db_encode=db.copy()
4
5 label_encoder = LabelEncoder()
6
7 db_encode['gender'] = label_encoder.fit_transform(db_encode['gender'])
8 db_encode['db_class'] = label_encoder.fit_transform(db_encode['db_class'])
9
10 db_encode.info()

```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  ---
 0   id          1000 non-null   int64
 1   no_patient  1000 non-null   int64
 2   gender      1000 non-null   int64
 3   age         1000 non-null   int64
 4   urea        1000 non-null   float64
 5   cr          1000 non-null   int64
 6   hba1c       1000 non-null   float64
 7   chol        1000 non-null   float64
 8   tg          1000 non-null   float64
 9   hdl         1000 non-null   float64
10   ldl         1000 non-null   float64
11   vldl        1000 non-null   float64
12   bmi         1000 non-null   float64
13   db_class    1000 non-null   int64
dtypes: float64(8), int64(6)
memory usage: 109.5 KB

```

```

1 #Compared the count with db_class count performed during cleaning to ensure consistency and understand the new format.
2 #Class Attributes are now: 0 = Not Diabetic (previously N) ; 1 = Pre-Diabetic (previously P) ; 2 = Diabetic (previously Y)
3
4 db_encode['db_class'].value_counts()

```




count	
db_class	
2	844
0	103
1	53

dtypes: int64

▼ Data Standardization

```

1 #Separating the original numerical and label encoded categorical variables to standardize the numerical variables only
2
3 num_cols = db_encode.drop(['gender', 'db_class'], axis=1).columns
4 cat_cols = ['gender', 'db_class']
5
6 std_scaler = StandardScaler()
7 db_encode[num_cols] = std_scaler.fit_transform(db_encode[num_cols])
8
9
10 #Combining the numerical and label encoded categorical variables
11
12 db_std = pd.concat([db_encode[num_cols], db_encode[cat_cols]], axis=1)
13 db_std.head(5)

```



	id	no_patient	age	urea	cr	hba1c	chol	tg	hdl	ldl	vldl	bmi	gender	db_class
0	0.672140	-0.074747	-0.401144	-0.144781	-0.382672	-1.334983	-0.509436	-1.035084	1.810756	-1.085457	-0.369958	-1.124622	0	0
1	1.641852	-0.069940	-3.130017	-0.212954	-0.115804	-1.334983	-0.893730	-0.678063	-0.158692	-0.457398	-0.342649	-1.326239	1	0
2	0.330868	-0.065869	-0.401144	-0.144781	-0.382672	-1.334983	-0.509436	-1.035084	1.810756	-1.085457	-0.369958	-1.124622	0	0
3	1.412950	-0.054126	-0.401144	-0.144781	-0.382672	-1.334983	-0.509436	-1.035084	1.810756	-1.085457	-0.369958	-1.124622	0	0
4	0.680463	-0.069939	-2.334096	0.673299	-0.382672	-1.334983	0.028576	-0.963680	-0.613180	-0.547121	-0.397267	-1.729472	1	0

Next steps:

[Generate code with db_std](#)[View recommended plots](#)[New interactive sheet](#)

▼ Data Splitting

```

1 #Splitting the standardized data into train and test set
2
3 train_std_set, test_std_set = train_test_split(db_std, train_size = 0.80, random_state = 89)
4
5
6 #Separating the feature variables and the target variable
7
8 X_train = train_std_set.drop('db_class', axis=1)
9 y_train = train_std_set['db_class']
10
11 X_test = test_std_set.drop('db_class', axis=1)
12 y_test = test_std_set['db_class']

```

▼ Recursive Feature Selection

```

1 #Creating a model to perform recursive feature selection
2
3 recursive_model = LogisticRegression(random_state=50)
4
5 rfe_selector = RFE(estimator=recursive_model, n_features_to_select=1)
6 rfe_selector.fit(X_train, y_train)
7
8
9 # Ranking and sorting the the features based on their relevance
10
11 feature_ranks = rfe_selector.ranking_
12
13 feature_rank_df = pd.DataFrame({'Feature': X_train.columns, 'Rank': feature_ranks})
14 feature_rank_df.columns = ['Feature', 'Rank']
15
16 feature_rank_df = feature_rank_df.sort_values(by='Rank')
17
18 print(feature_rank_df)

```

	Feature	Rank
5	hba1c	1
11	bmi	2
6	chol	3
12	gender	4
7	tg	5
10	vldl	6
2	age	7
3	urea	8
9	ldl	9
8	hdl	10
4	cr	11
0	id	12
1	no_patient	13

```

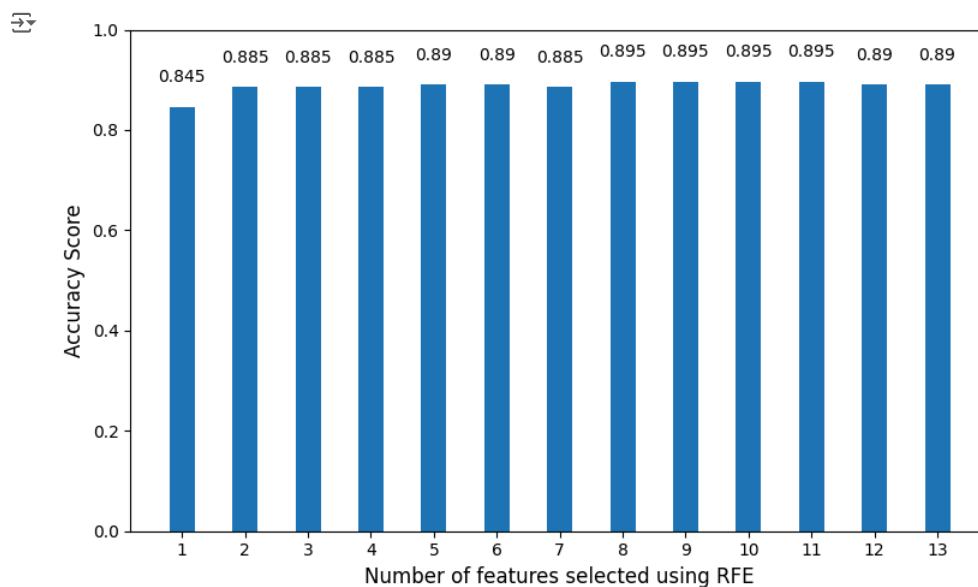
1 #Using a loop to iteratively perform recursive feature elimination, removing features one at a time, and calculating the accuracy score at each step
2
3 rfe_accuracy_list = []
4
5 for k in range(1,14):
6     rfe_selector = RFE(estimator=recursive_model, n_features_to_select=k)
7     rfe_selector.fit(X_train, y_train)
8
9     selected_features = X_train.columns[rfe_selector.support_]
10    X_train_rfe = X_train[selected_features]
11    X_test_rfe = X_test[selected_features]
12
13    recursive_model.fit(X_train_rfe, y_train)
14    rfe_preds = recursive_model.predict(X_test_rfe)
15
16    accuracy_rfe = round(accuracy_score(y_test, rfe_preds) ,3)
17
18    rfe_accuracy_list.append(accuracy_rfe)

```

```

1 #Creating chart to show the accuracy score as feature are progressively removed. This is used to determine the minimum number of features required for op
2
3 plt.figure(figsize=(8, 5))
4 plt.bar(np.arange(1,14), rfe_accuracy_list, width=0.4)
5
6 plt.xlabel('Number of features selected using RFE', fontsize=12)
7 plt.ylabel('Accuracy Score', fontsize=12)
8
9 plt.xticks(np.arange(1,14))
10 plt.ylim(0, 1.0)
11
12 for i, v in enumerate(rfe_accuracy_list):
13     plt.text(x=i+1, y=v+0.05, s=str(v), ha = 'center')
14
15 plt.tight_layout()

```



```

1 #Updating the number of features to select to 5 and showing the selected fatures
2
3 rfe_selector = RFE(estimator=recursive_model, n_features_to_select=8)
4 rfe_selector.fit(X_train, y_train)
5
6 selected_features = X_train.columns[rfe_selector.support_]
7 print("Selected Features:", selected_features)

```

```
1 #Reducing both the training and test dataset to the five selected features
2
3 X_train_rfe = X_train[selected_features]
4 X_test_rfe = X_test[selected_features]
```

- Base Models: Cross Validation and Predictions

```
1 #Applying StratifiedKFold for cross validation given the class imbalance in the dataset
2 #StratifiedKFold ensures that each fold maintains the same proportion of samples for each class and prevents misleading distribution in the validation fo
3
4 strat_fold = StratifiedKFold(n_splits=10, random_state=10, shuffle=True)
5
6 dc_model = DecisionTreeClassifier(max_depth=5, random_state=10, class_weight = 'balanced')
7 pred_dc_val = cross_val_predict(dc_model, X_train_rfe, y_train, cv=strat_fold)
8
9 lgr_model = LogisticRegression(random_state=10, multi_class='multinomial', class_weight = 'balanced')
10 pred_lgr_val = cross_val_predict(lgr_model, X_train_rfe, y_train, cv=strat_fold)
11
12 knn_model = KNeighborsClassifier(n_neighbors=5, weights='distance')
13 pred_knn_val = cross_val_predict(knn_model, X_train_rfe, y_train, cv=strat_fold)
```

[illegible]

- Second Level Model: Stacking and Training

```
1 # Stacking the validation predictions from each base model to form the training set for the second level model
2
3 X_seclvl_train = np.column_stack((pred_dc_val, pred_lgr_val, pred_knn_val))
```

```
1 # Training the second level model, Logistic Regression, with the stacked predictions
2
3 seclvl_model = LogisticRegression(random_state=20, multi_class='multinomial', class_weight = 'balanced')
4 seclvl_model.fit(X_seclvl_train, y_train)
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:1247: FutureWarning: 'multi_class' was deprecated in version 1.5 and will be removed in a future version.
  warnings.warn(

```

```
LogisticRegression
LogisticRegression(class_weight='balanced', multi_class='multinomial',
                    random_state=20)
```

- Final Prediction

```
1 #Training the base models on all of the training dataset to adequately make predictions on the test set
2
3 dc_model.fit(X_train_rfe, y_train)
4 pred_dc_test = dc_model.predict(X_test_rfe)
5
6 lgr_model.fit(X_train_rfe, y_train)
7 pred_lgr_test = lgr_model.predict(X_test_rfe)
8
9 knn_model.fit(X_train_rfe, y_train)
10 pred_knn_test = knn_model.predict(X_test_rfe)
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:1247: FutureWarning: 'multi_class' was deprecated in version 1.5 and will be removed in version 1.7:
  warnings.warn(

```

```

1 # Stacking the predictions made by the base models to create a new testing dataset for the second level model
2 X_seclvl_test = np.column_stack((pred_dc_test, pred_lgr_test, pred_knn_test))
3
4 # Final prediction made using the second level model
5 pred_final = seclvl_model.predict(X_seclvl_test)

```

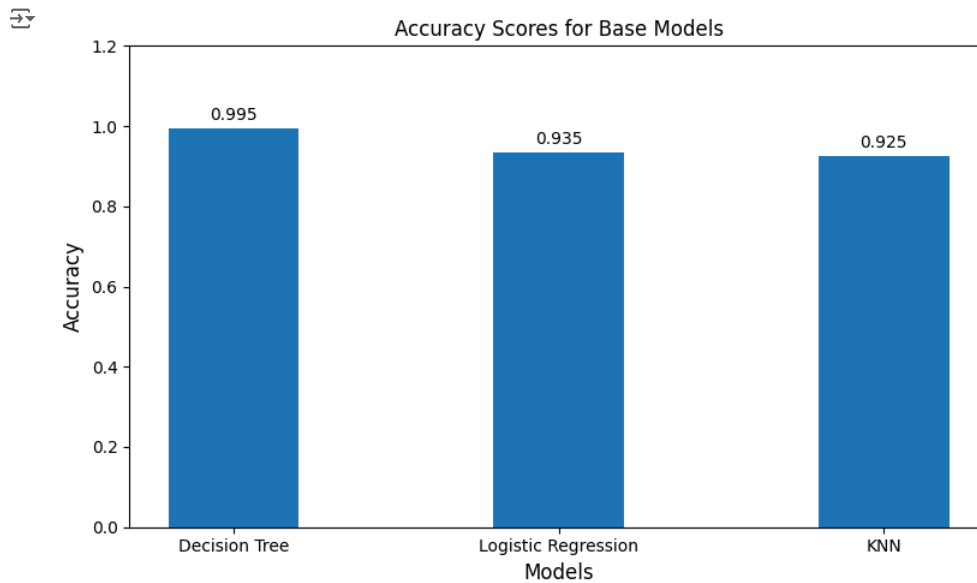
Stage 5: Model Evaluation

Base Models Evaluation

```

1 #Creating barchart to show the accuracy score among the base models
2
3 accuracy_dc = accuracy_score(y_test, pred_dc_test)
4 accuracy_lgr = accuracy_score(y_test, pred_lgr_test)
5 accuracy_knn = accuracy_score(y_test, pred_knn_test)
6
7 accuracy_base_scores = [accuracy_dc, accuracy_lgr, accuracy_knn]
8 model_names_final = ['Decision Tree', 'Logistic Regression', 'KNN']
9
10 plt.figure(figsize=(8, 5))
11 plt.bar(model_names_final, accuracy_base_scores, width=0.4)
12
13 plt.xlabel('Models', fontsize=12)
14 plt.ylabel('Accuracy', fontsize=12)
15 plt.title('Accuracy Scores for Base Models', fontsize=12)
16 plt.ylim(0, 1.2)
17
18 for i, v in enumerate(accuracy_base_scores):
19     plt.text(i, v + 0.02, f"{v:.3f}", ha='center', fontsize=10)
20
21 plt.tight_layout()
22 plt.show()

```



```

1 cf_dc = confusion_matrix(y_test, pred_dc_test)
2 print("Confusion Matrix for Decision Tree Base Model:\n")
3 print(cf_dc)
4
5 print('\n')
6
7 print("Performance Metrics for Decision Tree Base Model:\n")
8 print(classification_report(y_test, pred_dc_test))

```

Confusion Matrix for Decision Tree Base Model:

```

[[ 19  0  0]
 [ 0 16  1]
 [ 0  0 164]]

```

Performance Metrics for Decision Tree Base Model:

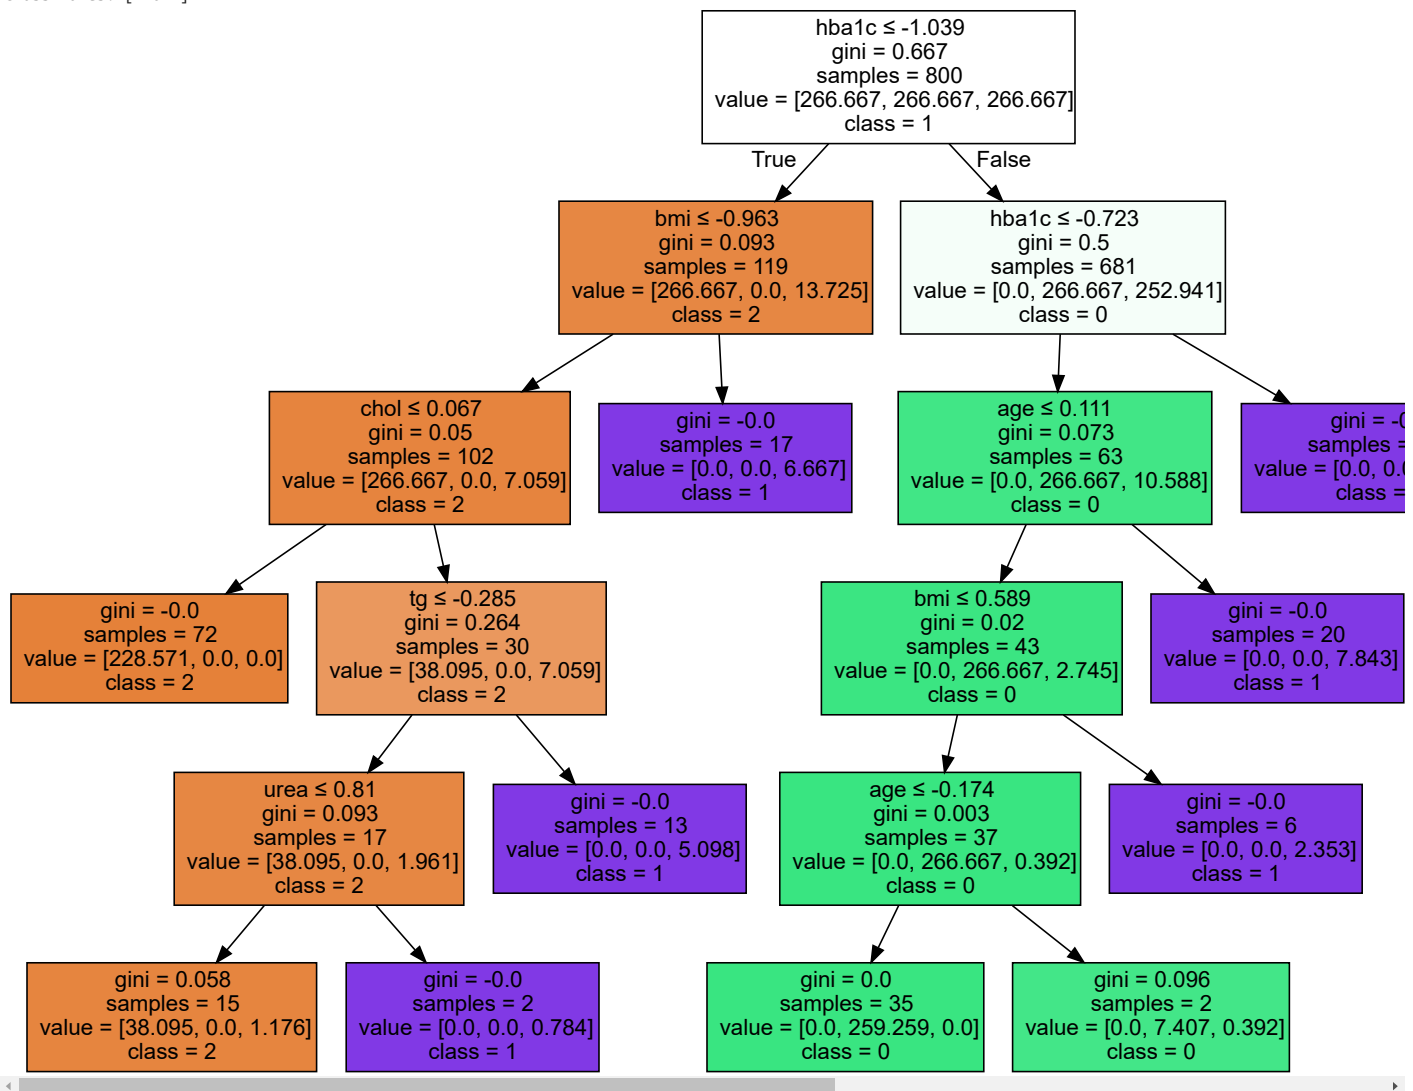
	precision	recall	f1-score	support
0	1.00	1.00	1.00	19
1	1.00	0.94	0.97	17
2	0.99	1.00	1.00	164
accuracy			0.99	200
macro avg	1.00	0.98	0.99	200
weighted avg	1.00	0.99	0.99	200

```

1 class_values = y_train.unique() # Or, if it's in a DataFrame, use df['Class'].unique()
2 print("Class Names:", class_values)
3
4 # Visualize the Decision Tree
5 dot_data = tree.export_graphviz(dc_model, out_file=None,
6                                 feature_names=X_train_rfe.columns, # Use the column names for features
7                                 class_names=[str(c) for c in class_values], # Class names as strings
8                                 filled=True, # Color the nodes based on their class
9                                 special_characters=True) # Handle special characters in feature names
10
11 # Generate and render the graph
12 graph = graphviz.Source(dot_data, format="png")
13 graph

```

Class Names: [2 0 1]



```

1 cf_lgr = confusion_matrix(y_test, pred_lgr_test)
2 print("Confusion Matrix for Logisitic Regression Base Model:\n")
3 print(cf_lgr)
4
5 print('\n')
6
7 print("Performance Metrics for Logisitic Regression Base Model:\n")
8 print(classification_report(y_test, pred_lgr_test))

```

↗ Confusion Matrix for Logisitic Regression Base Model:

```
[[ 17  2  0]
 [  0 17  0]
 [  4  7 153]]
```

Performance Metrics for Logisitic Regression Base Model:

	precision	recall	f1-score	support
0	0.81	0.89	0.85	19
1	0.65	1.00	0.79	17
2	1.00	0.93	0.97	164
accuracy			0.94	200
macro avg	0.82	0.94	0.87	200
weighted avg	0.95	0.94	0.94	200

```
1 cf_knn =confusion_matrix(y_test, pred_knn_test)
2 print("Confusion Matrix for kNN Base Model:\n")
3 print(cf_knn)
4
5 print('\n')
6
7 print("Performance Metrics for kNN Base Model:\n")
8 print(classification_report(y_test, pred_knn_test))
```

↗ Confusion Matrix for kNN Base Model:


```
[[ 15  1  3]
 [  2 12  3]
 [  2  4 158]]
```

Performance Metrics for kNN Base Model:

	precision	recall	f1-score	support
0	0.79	0.79	0.79	19
1	0.71	0.71	0.71	17
2	0.96	0.96	0.96	164
accuracy			0.93	200
macro avg	0.82	0.82	0.82	200
weighted avg	0.93	0.93	0.93	200

▼ Base Model Performance with Cross Validation

```
1 #Creating barchart to show the accuracy score among the cross validated base models
2
3 acc_dc_val = accuracy_score(y_train, pred_dc_val)
4 acc_lgr_val = accuracy_score(y_train, pred_lgr_val)
5 acc_knn_val = accuracy_score(y_train, pred_knn_val)
6
7 acc_val = [acc_dc_val, acc_lgr_val, acc_knn_val]
8 model_names_val = ['Decision Tree', 'Logistic Regression', 'KNN']
9
10 plt.figure(figsize=(8, 5))
11 plt.bar(model_names_val, acc_val, width=0.4)
12
13 plt.xlabel('Models', fontsize=12)
14 plt.ylabel('Accuracy', fontsize=12)
15 plt.title('Accuaracy Scores for Base Models with Cross Validation', fontsize=12)
16
17 for i, v in enumerate(acc_val):
18     plt.text(i, v + 0.02, f"{v:.3f}", ha='center', fontsize=10)
19
20 plt.tight_layout()
21 plt.show()
```



```
1 cf_dc_val = confusion_matrix(y_train, pred_dc_val)
2 print("Confusion Matrix for Decision Tree Base Model with Cross Validation:\n")
3 print(cf_dc_val)
4
5 print('\n')
6
7 print("Performance Metrics for Decision Tree Base Model with Cross Validation:\n")
8 print(classification_report(y_train, pred_dc_val))
```

Confusion Matrix for Decision Tree Base Model with Cross Validation:

```
[[ 81  0  3]
 [ 0 34  2]
 [ 4  1 675]]
```

Performance Metrics for Decision Tree Base Model with Cross Validation:

	precision	recall	f1-score	support
0	0.95	0.96	0.96	84
1	0.97	0.94	0.96	36
2	0.99	0.99	0.99	680
accuracy			0.99	800
macro avg	0.97	0.97	0.97	800
weighted avg	0.99	0.99	0.99	800

```
1 cf_lgr_val =confusion_matrix(y_train, pred_lgr_val)
2 print("Confusion Matrix for Logisitic Regression Base Model with Cross Validation:\n")
3 print(cf_lgr_val)
4
5 print('\n')
6
7 print("Performance Metrics for Logisitic Regression Base Model with Cross Validation:\n")
8 print(classification_report(y_train, pred_lgr_val))
```

Confusion Matrix for Logisitic Regression Base Model with Cross Validation:

```
[[ 74 10  0]
 [ 3 31  2]
 [23 34 623]]
```

Performance Metrics for Logisitic Regression Base Model with Cross Validation:

	precision	recall	f1-score	support
0	0.74	0.88	0.80	84
1	0.41	0.86	0.56	36
2	1.00	0.92	0.95	680
accuracy			0.91	800
macro avg	0.72	0.89	0.77	800
weighted avg	0.94	0.91	0.92	800

```
1 cf_knn_val =confusion_matrix(y_train, pred_knn_val)
2 print("Confusion Matrix for kNN Base Model with Cross Validation:\n")
```

```
3 print(cf_knn_val)
4
5 print('\n')
6
7 print("Performance Metrics for kNN Base Model with Cross Validation:\n")
8 print(classification_report(v_train, pred_knn_val))
```

Confusion Matrix for kNN Base Model with Cross Validation:

```
[[ 62   9  13]
 [  7  24   5]
 [  8   6 666]]
```

Performance Metrics for kNN Base Model with Cross Validation:

	precision	recall	f1-score	support
0	0.81	0.74	0.77	84
1	0.62	0.67	0.64	36
2	0.97	0.98	0.98	680
accuracy			0.94	800
macro avg	0.80	0.79	0.80	800
weighted avg	0.94	0.94	0.94	800

Second Level Model Evaluation

```
1 #Creating barchart to compare the accuracy scores for the base and second level models
2
3 accuracy_seclvl = accuracy_score(y_test, pred_final)
4
5 accuracy_scores = [acc_dc_val , acc_lgr_val, acc_knn_val, accuracy_seclvl ]
6 model_names_final = ['Decision Tree', 'Logistic Regression', 'KNN', 'Second-Level Model']
7
8
9 plt.figure(figsize=(8, 5))
10 plt.bar(model_names_final, accuracy_scores, width=0.4)
11
12 plt.xlabel('Models', fontsize=12)
13 plt.ylabel('Accuracy', fontsize=12)
14 plt.title('Accuracy Scores for Cross Validated Base Models and Second Level Model', fontsize=12)
15 plt.ylim(0, 1.2)
16
17 for i, v in enumerate(accuracy_scores):
18     plt.text(i, v + 0.02, f"{v:.3f}", ha='center', fontsize=10)
19
20 plt.tight_layout()
21 plt.show()
```

